



deeplearning.ai


Convolutional Neural Networks

Why convolutions?

Why convolutions

2 Adv over FC layers

- ① parameter sharing
- ② sparsity of connections



Actual params

$$\begin{aligned} 5 \times 5 &\rightarrow \text{filter} \quad \# \text{ filters} \\ 25 \times 6 &= 150 \\ &+ 6 \text{ bias params} \\ &= 156 \text{ params} \end{aligned}$$

Imagine if these 2 layers were fully connected, then

$$\begin{aligned} \# \text{ params} &= n^{[l+1]} \times n^{[l]} \\ &= 4704 \times 3072 \\ &\approx 12 \text{ M (lot of params for a small image)} \end{aligned}$$

Why convolutions

10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0
10	10	10	0	0	0

*

1	0	-1
1	0	-1
1	0	-1

=

0	30	30	0
0	30	30	0
0	30	30	0
0	30	30	0

Note
CNNs also are "translation Invariant"
meaning If Image of cat shifts to the
right, it still is able to detect a cat,
why? Because the filters are all applied
to each iteration of the stride & the
shifted Images' "cat features" will be
captured in a later/earlier strides
convolution


← This came
from a
very small
of calculations
ie, only 9 cells
were involved
⇒ this cell is
"derived from"
or "connected to"
only 9 (out of 36)
cells ⇒ sparse

Parameter sharing: A feature detector (such as a vertical edge detector) that's useful in one part of the image is probably useful in another part of the image. Same filter is used across all stride iterations of the Image to detect vertical edge

Sparsity of connections: In each layer, each output value depends only on a small number of inputs.

Putting it together

Training set $(x^{(1)}, y^{(1)}) \dots (x^{(m)}, y^{(m)})$.



$$\text{Cost } J = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)})$$

Use gradient descent to optimize parameters to reduce J